

# Best Practices for Stream Processing with GridGain and Apache Ignite and Kafka



Alexey Kukushkin  
Professional Services



Rob Meyer  
Outbound Product Management

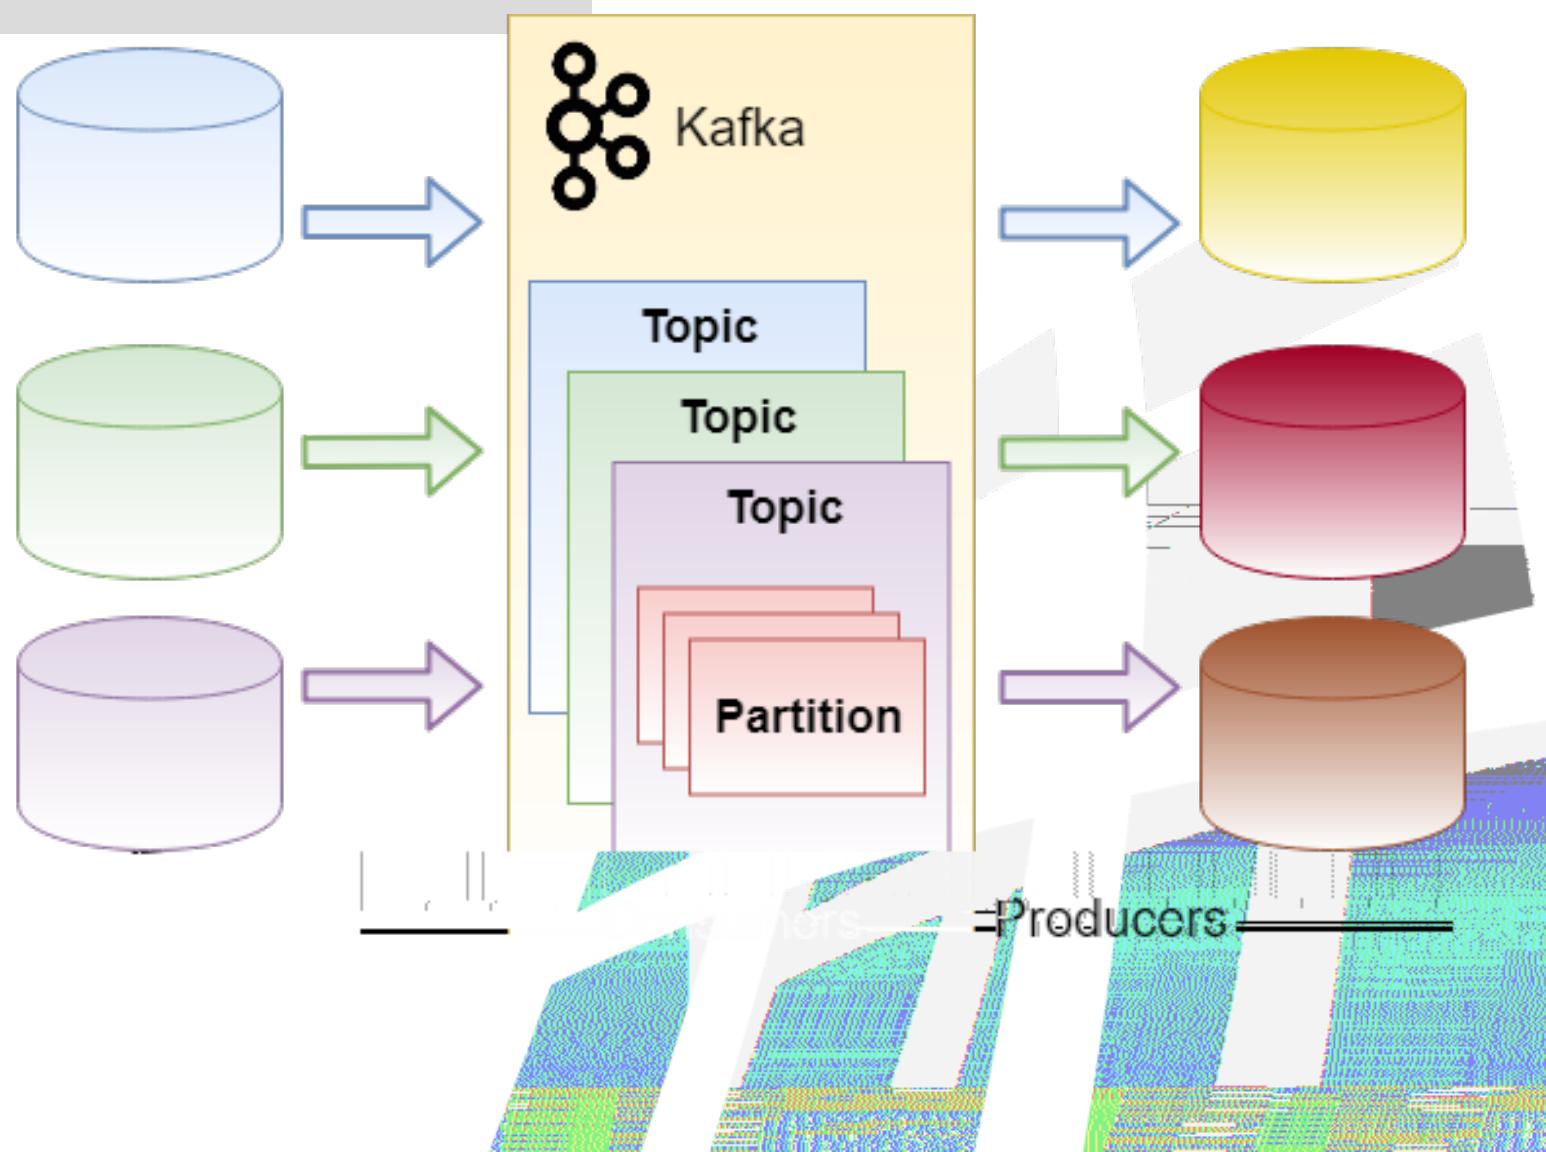
# Agenda

- Why we need Kafka/Confluent-Ignite/GridGain integration
- Ignite/GridGain Kafka/Confluent Connectors
- Deployment, monitoring and management
- Integration Examples
- Performance and scalability tuning
- Q & A

# Why we need Kafka/Confluent-Ignite/GridGain integration

# Apache Kafka and Confluent

- A distributed streaming platform:
- Publish/subscribe
  - Scalable
  - Fault-tolerant
  - Real-time
  - Persistent
  - Written mostly in Scala



# GridGain In-Memory Computing Platform

- Built on Apache Ignite
  - Comprehensive platform that supports all projects
  - No rip and replace
  - In-memory **speed**, petabyte **scale**
  - Enables HTAP, streaming analytics and continuous learning
- What GridGain adds
  - Production-ready releases
  - Enterprise-grade security, deployment and management
  - Global support and services
  - Proven for mission critical apps



Existing Applications



New Applications, Analytics



Streaming, Machine Learning

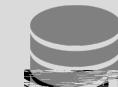
In-Memory  
Data Grid

In-Memory  
Database

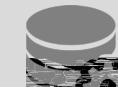
Streaming  
Analytics

Continuous  
Learning

GridGain In-Memory Computing Platform



RDBMS

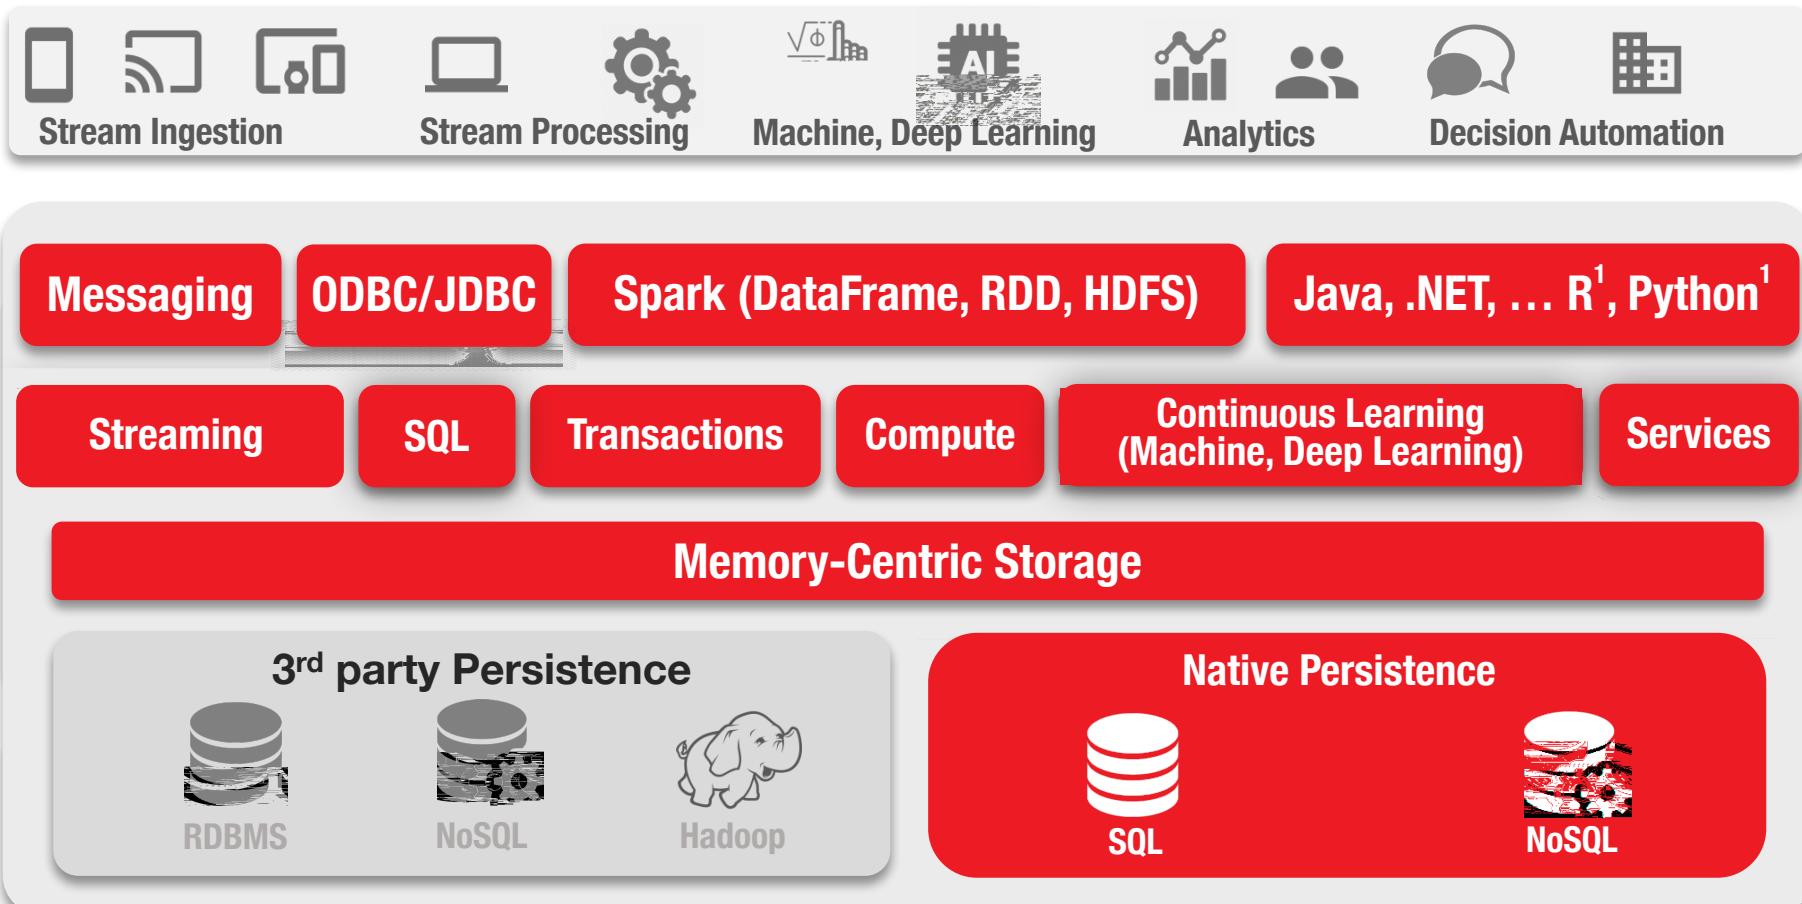


NoSQL



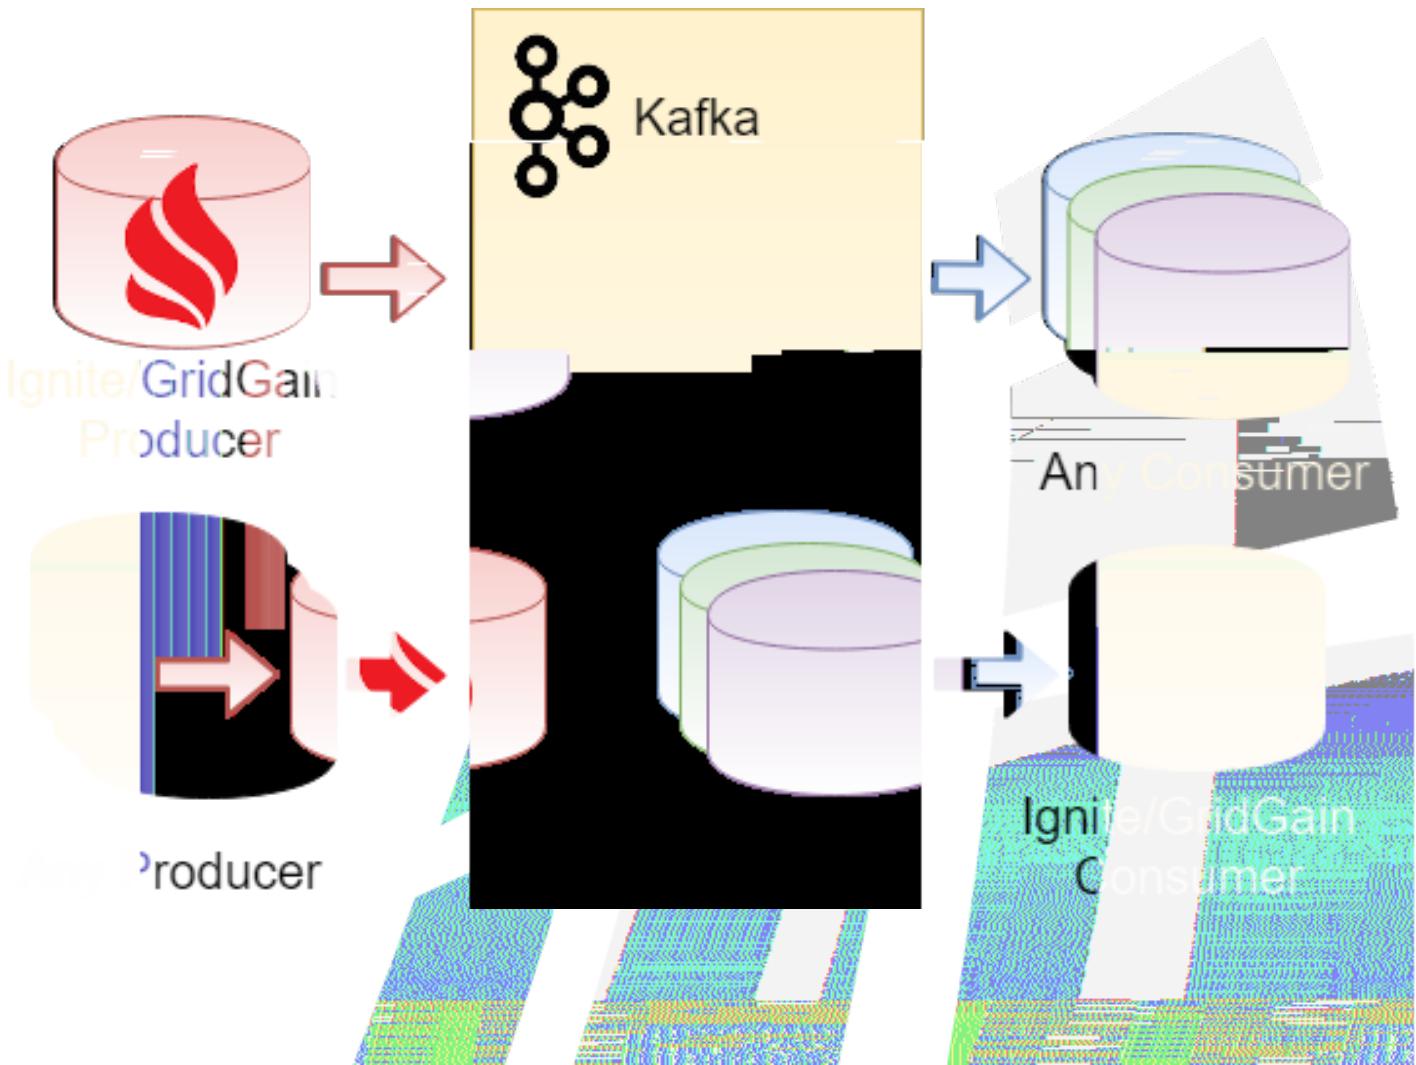
Hadoop

# Streaming Analytics, Machine and Deep Learning

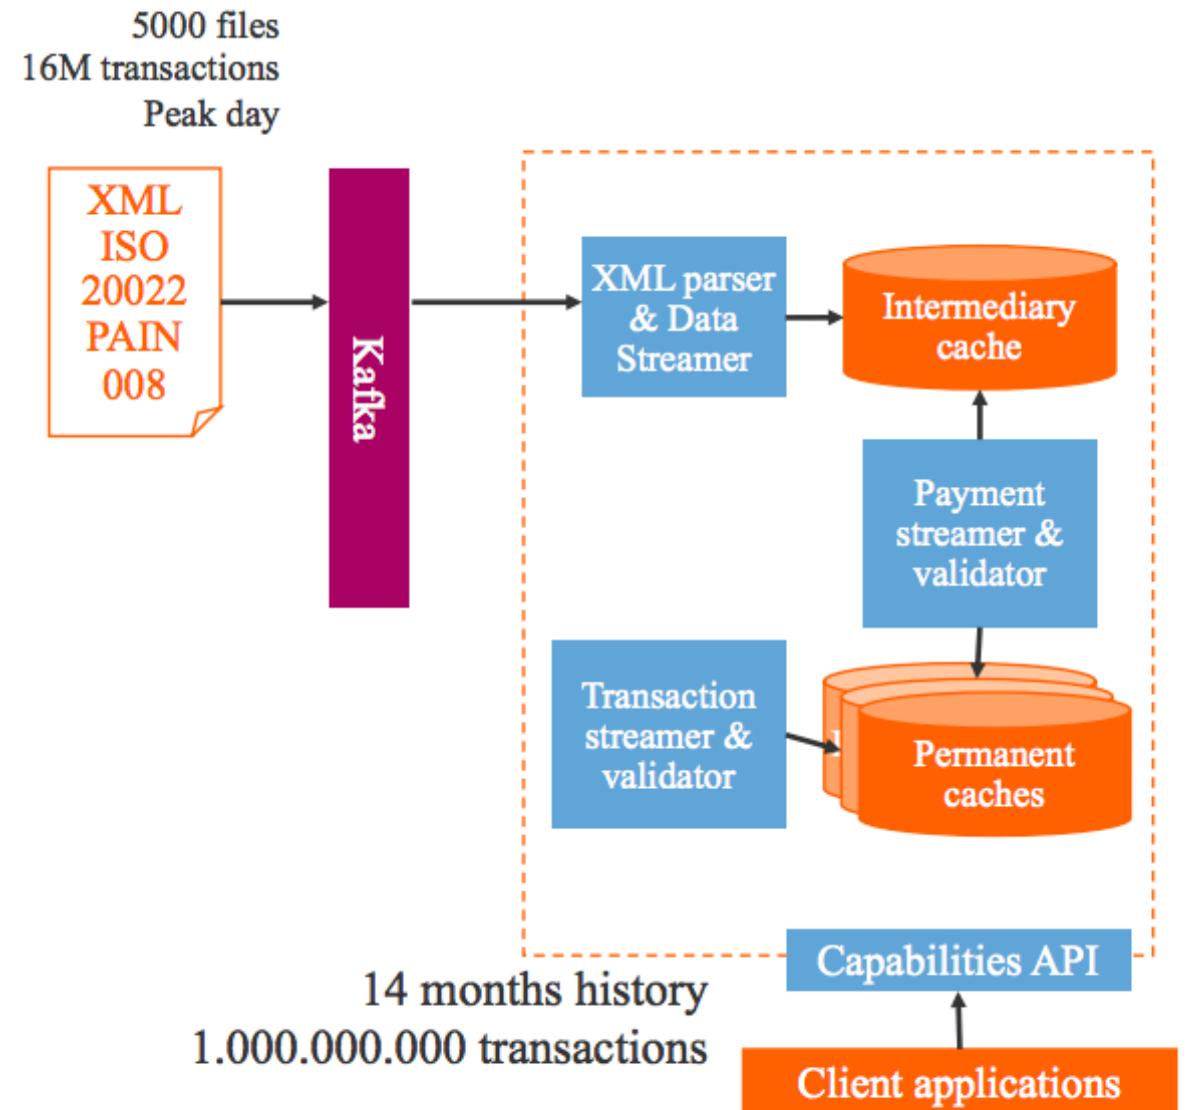
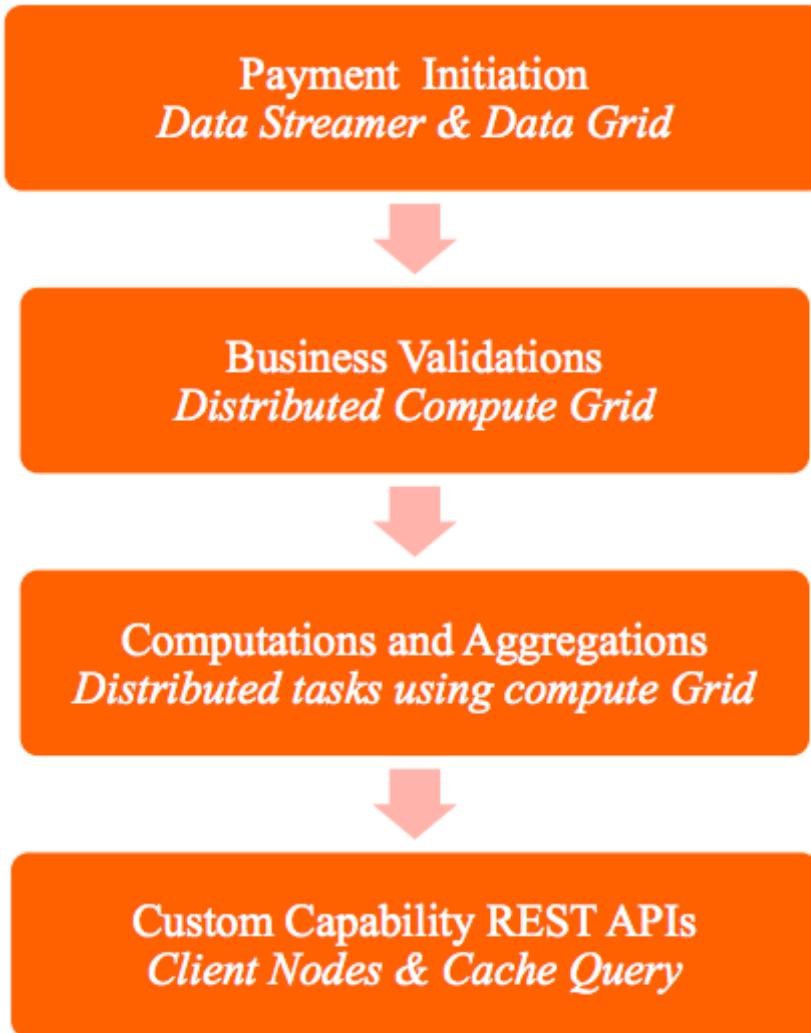


# Ignite/GridGain & Kafka Integration

- Kafka is commonly used as a messaging backbone in a heterogeneous system
- Add Ignite/GridGain to a Kafka-based system



## SEPA DD



# ignite and GridGain Kafka Connectors



# Developing Kafka Consumers & Producers

You can develop Kafka integration for any system using Kafka Producer and Consumer APIs but you need to solve problems like:

- How to use each API of every producer and consumer
- How Kafka will understand your data
- How data will be converted between producers and consumers
- How to scale the producer-to-consumer flow
- How to recover from a failure
- ... and many more

# GridGain-Kafka Connector: Out-of-the-box Integration

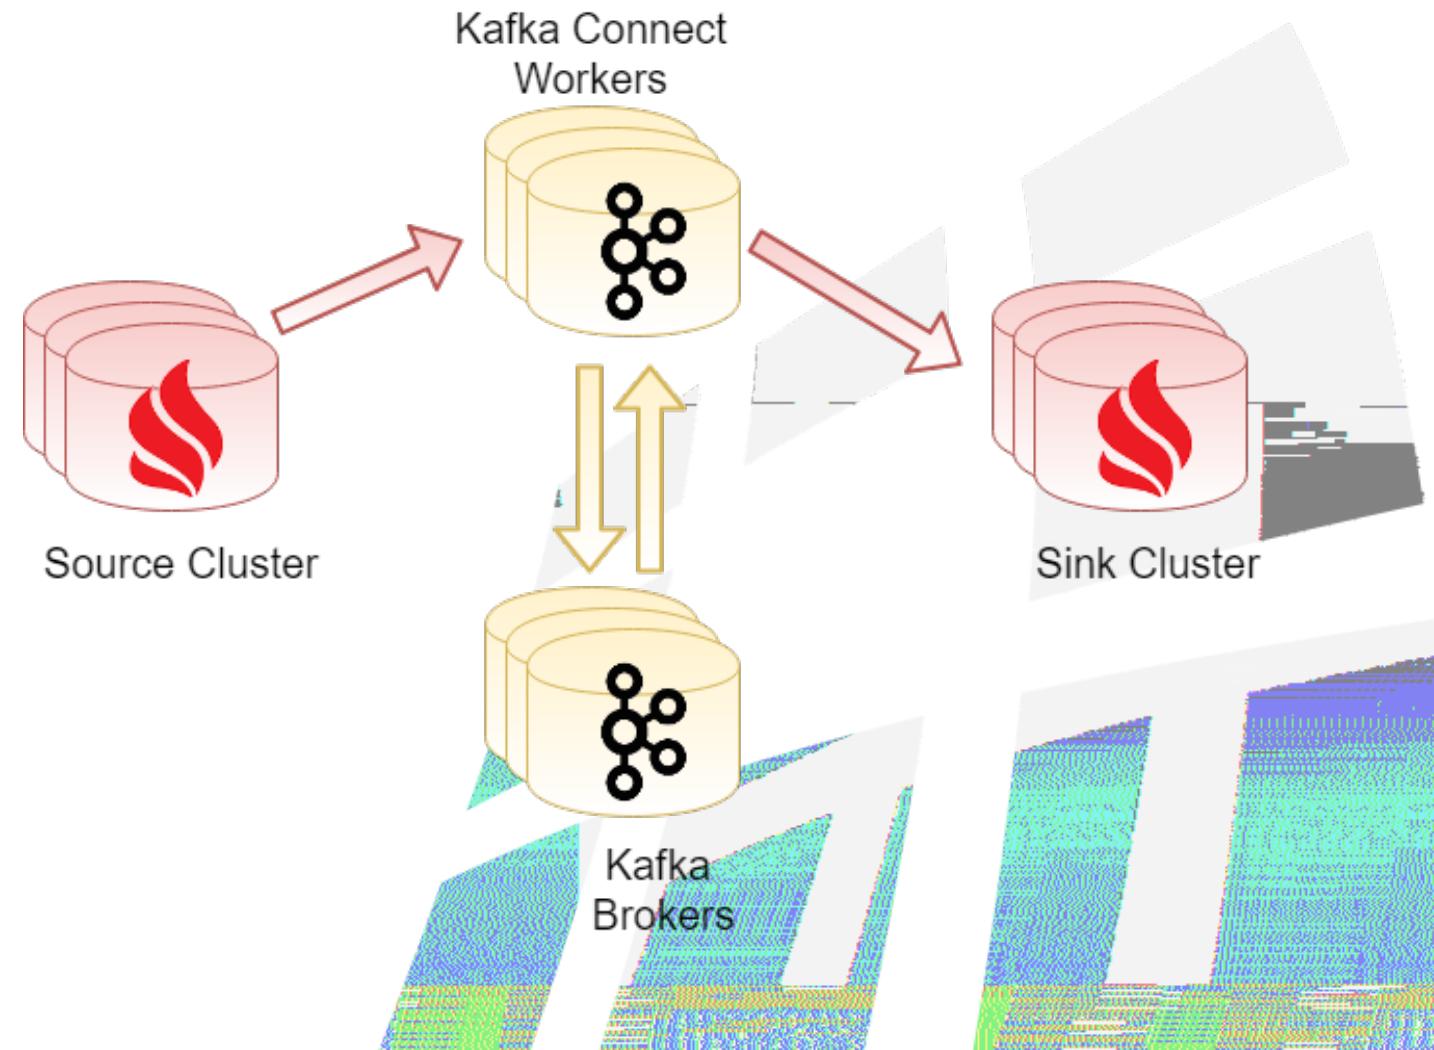
- Addresses all the integration challenges using best practices
- Does not need any coding even in the most complex integrations
- Developed by GridGain/Ignite Community with help from Confluent to ensure both Ignite and Kafka best practices
- Based on Kafka Connect and Ignite APIs
  - Kafka Connect API encourages design for scalability, failover and data schema
  - GridGain Source Connector uses Ignite Continuous Queries
  - GridGain Sink Connector uses Ignite Data Streamer



# Kafka Connect Server Types

In general, there are 4 separate clusters in Kafka Connect infrastructure:

- Kafka cluster
  - cluster nodes called **Brokers**
- Kafka Connect cluster
  - cluster nodes called **Workers**
- Source and Sink GridGain/Ignite clusters
  - **Server Nodes**



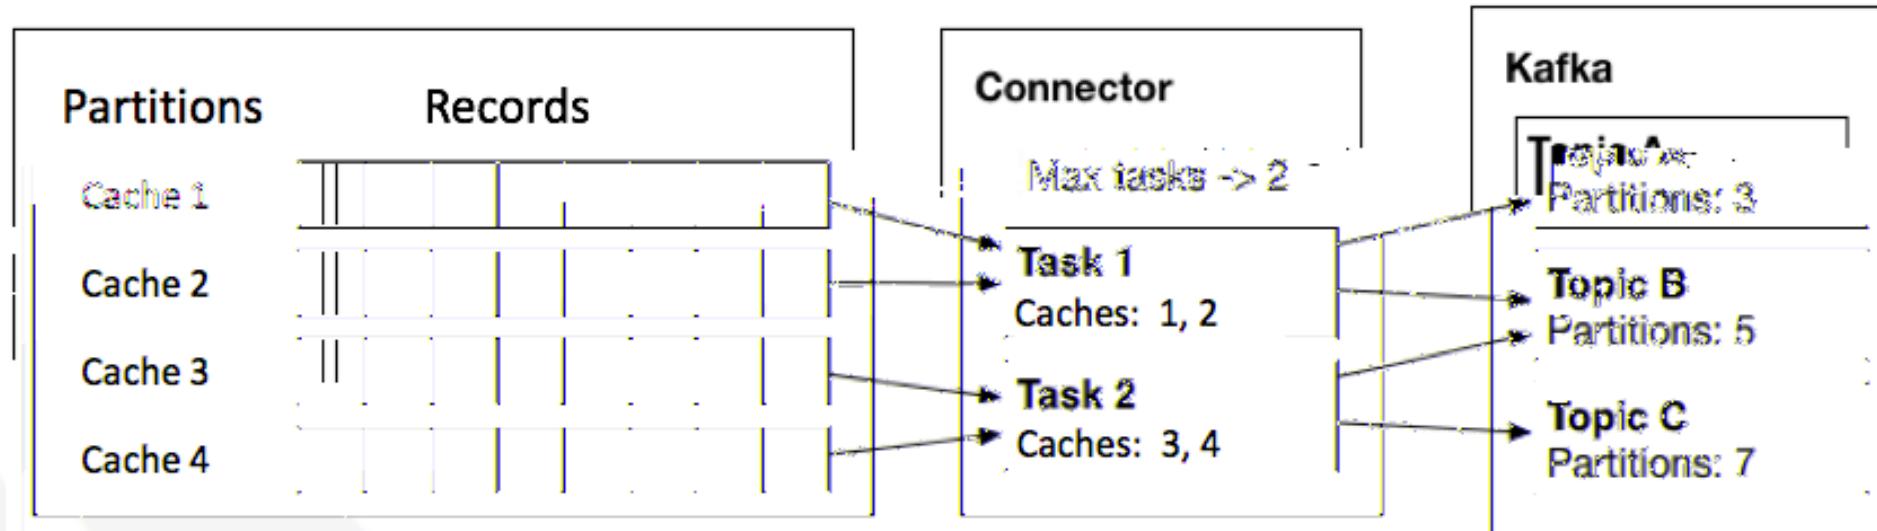
# GridGain Connector Features

Two connectors independent from each other:

- GridGain Source Connector
  - streams data from GridGain into Kafka
  - uses Ignite continuous queries
- GridGain Sink Connector
  - streams data from Kafka into GridGain
  - uses Ignite data streamers

# GridGain Source Connector: Scalability

Kafka Source Connector Model:



Scales by assigning multiple source **partitions** to Kafka Connect **tasks**.

For GridGain Source Connector:

- Partition = Cache
- Record = Cache Entry

# GridGain Source Connector: Rebalancing and Failover

**Rebalancing:** re-assignment of Kafka Connectors and Tasks to Workers when

- A Worker joins or leaves the cluster
- A cache is added or removed

**Failover:** resuming operation after a failure

- how to resume after failure or rebalancing without losing cache updates occurred when Kafka Worker node was down?

**Source Offset:** position in the source stream. Kafka Connect:

- provides persistent and distributed source offset storage
- automatically saves last committed offset
- allows resuming from the last offset without losing data.

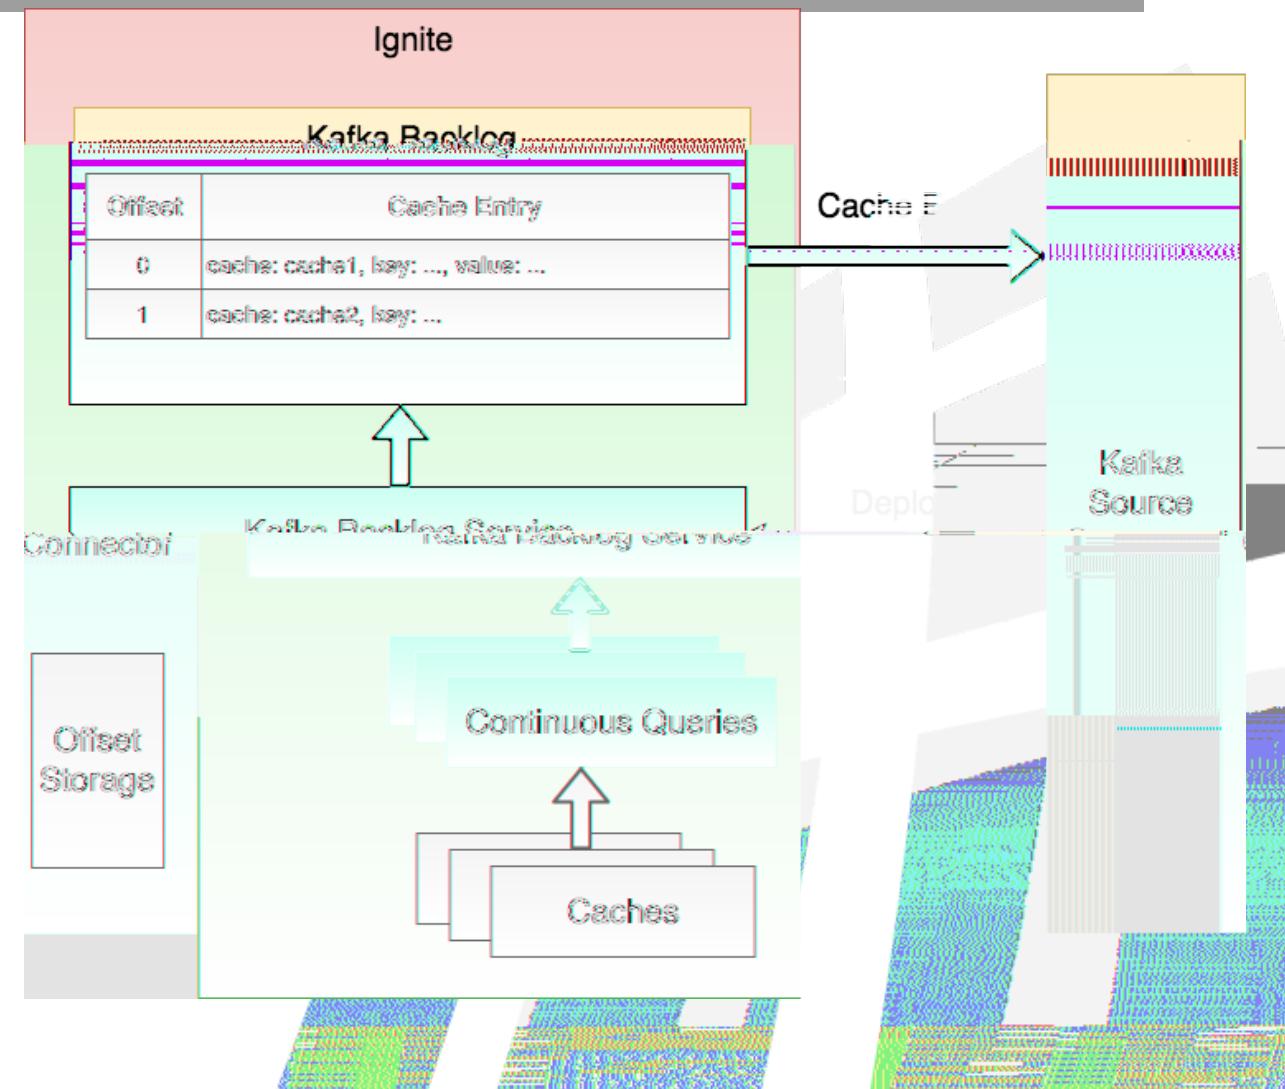
**Problem:** caches have no offsets!

# GridGain Source Connector: Failover Policies

- **None:** no source offset saved, start listening to current data after restart
  - Cons: updates occurred during downtime are lost (“at least once” data delivery guarantee violated)
  - Pros: fastest
- **Full Snapshot:** no source offset saved, always pull all data from the cache upon startup
  - Cons:
    - Slow, not applicable for big caches
    - Duplicate data (“exactly once” data delivery guarantee is violated)
  - Pros: no data is lost

# GridGain Source Connector: Failover Policies

- **Backlog:** resume from the last committed source offset
  - Kafka Backlog cache in Ignite
    - Key: incremental offset
    - Value: cache name and serialized cache entries
  - Kafka Backlog service in Ignite
    - Runs continuous queries pulling data from source caches into Backlog
  - Source Connector gets data from Backlog from backlog starting from the last committed offset
  - Cons
    - Intrusive: GridGain cluster impact
    - Complex configuration: need to estimate amount of memory for Backlog



# GridGain Source Connector: Dynamic Reconfiguration

Connector monitors list of available caches and re-configures itself if a cache is added or removed.

Use `gridGainCacheNames` and  
caches to pull data.

`gridGainCacheNames` properties to define from which

# GridGain Source Connector: Initial Data Load

Use `initialDataLoadEnabled` configuration property to specify if you want the Connector to load the data that is already in the cache by the time the Connector starts.

# GridGain Sink Connector

- Sink Connectors are inherently scalable since consuming data from a Kafka topic is scalable
- Sink Connectors inherently support failover thanks to the Kafka Connector framework auto-committing offsets of the pushed data.

# GridGain Connector Data Schema

Both Source and Sink GridGain Connectors support data schema.

- Allows GridGain Connectors understand data with attached schema from other Kafka producers and consumers
- Source Connector attaches Kafka schema built from Ignite Binary objects
- Sink Connector converts Kafka records to Ignite Binary objects using Kafka schema

Limitations:

- Ignite Annotations are not supported
- Ignite CHAR converted to Kafka SHORT (same for arrays)
- Ignite UUID and CLASS converted to Kafka STRING (same for arrays)

# Ignite Connector Features

- Ignite Source Connector
  - pushes data from Ignite into Kafka
  - uses Ignite Events
    - must enable `EVT_CACHE_OBJECT_PUT`, which negatively impacts cluster performance
- Ignite Sink Connector
  - pulls data from Kafka into Ignite
  - use Ignite data streamer

# Apache Ignite vs. GridGain Connectors

Feature	Apache Ignite Connector	GridGain Connector
Scalability	<b>Limited</b> Source connector is not parallel Sink connector is parallel	Source connector creates a task per cache Sink connector is parallel
Failover	<b>NO</b> Source data is lost during connector restart or rebalancing	<b>YES</b> Source connector can be configured to resume from the last committed offset
Preserving source data schema	<b>NO</b>	<b>YES</b>
Handling multiple caches	<b>NO</b>	<b>YES</b> Connector can be configured to handle any number of caches
Dynamic Reconfiguration	<b>NO</b>	<b>YES</b> Source connector detects added or removed caches and re-configures itself

# Apache Ignite vs. GridGain Connectors

Feature	Apache Ignite Connector	GridGain Connector
Initial Data Load	<b>NO</b>	<b>YES</b>
Handling data removals	<b>YES</b>	<b>YES</b>
Serialization and Deserialization of data	<b>YES</b>	<b>YES</b>
Filtering	<b>Limited</b> Only source connector supports a filter	<b>YES</b> Both source and sink connectors support filters
Transformations	Kafka SMTs	Kafka SMTs

# Apache Ignite vs. GridGain Connectors

Feature	Apache Ignite Connector	GridGain Connector
DevOps	Some free-text error logging	Health Model defined
Support	Apache Ignite Community	Supported by GridGain, certified by Confluent
Packaging	Uber JAR	Connector Package
Deployment	Plugin PATH on all Kafka Connect workers	Plugin PATH on all Kafka Connect workers. CLASSPATH on all GridGain nodes.
Kafka API Version	0.10	2.0
Source API	Ignite events	Ignite continuous queries
Sink API	Ignite data streamer	Ignite data streamer



# **Deployment, monitoring and management**

# GridGain Connector Deployment

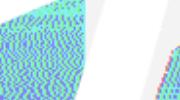
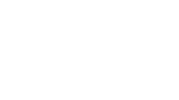
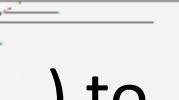
1. Prepare Connector Package
2. Register Connector with Kafka
3. Register Connector with GridGain

# Prepare GridGain Connector Package

1. GridGain-Kafka Connector is part of GridGain Enterprise and Ultimate 8.5.3 (to be released in the end of October 2018)
2. The connector is in
  - (`GRIDGAIN_HOME` environment variable points to the root GridGain installation directory)
3. Pull missing connector dependencies into the package:

# Register GridGain Connector with Kafka

For every Kafka Connect Worker:

1. Copy GridGain Connector package directory to where you want Kafka Connectors to be located  
for example, into  directory
2. Edit Kafka Connect Worker configuration (  or  ) to register the connector on the plugin path:  
  
  
  
  
  
  
  
  
  
  
  
  
  
  


# Register GridGain Connector with GridGain

This assumes GridGain version is 8.5.3

On every GridGain server node copy the below JARs into  
directory. Get the Kafka JARs from the Kafka

Connect workers:

- 
- 
-

# Ignite Connector Deployment

1. Prepare Connector Package
2. Register Connector with Kafka

# Prepare Ignite Connector Package

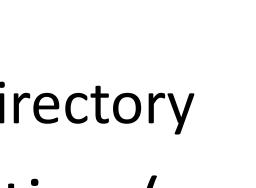
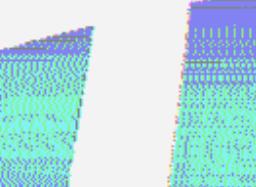
This assumes Ignite version is 2.6.

Create a directory containing the below JARs (find JARs in the sub-directories):

- •  
•  
•  
•  
•  
•  
•  
•  
•

# Register GridGain Connector with Kafka

For every Kafka Connect Worker:

1. Copy Ignite Connector package directory to where you want Kafka Connectors to be located  
for example, into  directory
2. Edit Kafka Connect Worker configuration (  or  ) to register the connector on the plugin path:  
  
  


# Monitoring: GridGain Connector

Well defined Health Model:

- Numeric Event ID uniquely identifies specific problem
- Event severity
- Problem description and recovery action is available at  
<https://docs.gridgain.com/docs/certified-kafka-connector-monitoring>

Configure your monitoring system to detect event ID in the logs and may be run automated recovery as defined in the Health Model

- Sample structured log entry (# used as a delimiter):

# Monitoring: Ignite Connector

No Health Model is defined.

1. Run negative tests
2. Check Kafka and Ignite logs output
3. Configure your monitoring system to detect corresponding text patterns in the logs

# Integration Examples

Propagating RDBMS updates into GridGain

# Propagating RDBMS updates into GridGain

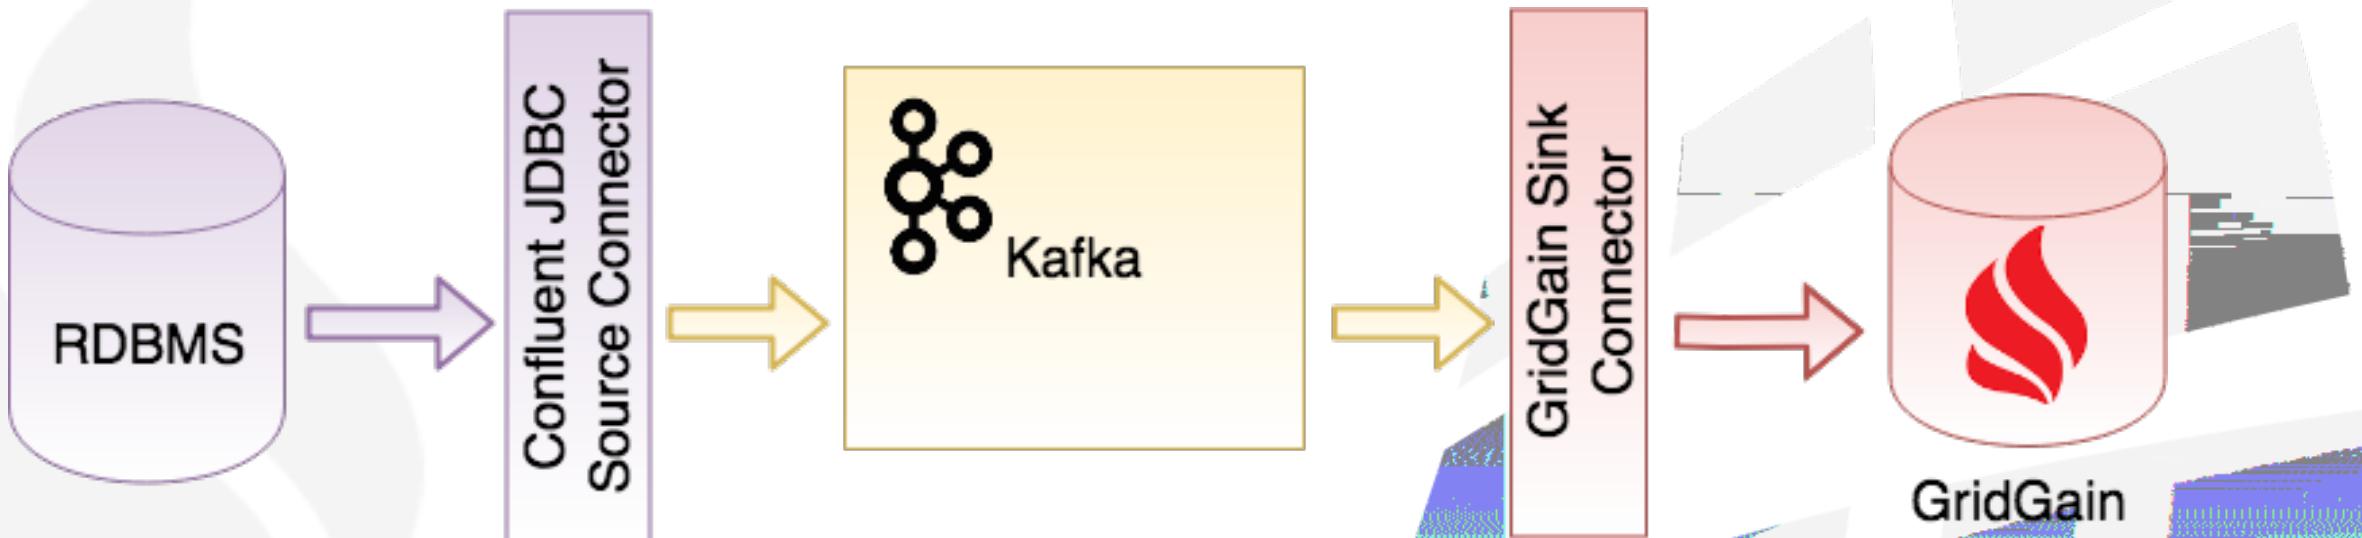
Ignite/GridGain has a 3<sup>rd</sup> Party Persistence feature (Cache Store) that allows:

- Propagating cache changes to external storage like RDBMS
- Automatically copying data from external storage to Ignite upon accessing data missed in Ignite

What if you want to propagate external storage change to Ignite at the moment of the change? - 3<sup>rd</sup> Party Persistence cannot do that!

# Propagating RDBMS updates into GridGain

Use Kafka to achieve that without writing single line of code!



# Assumptions

- For simplicity we will run everything on the same host
  - In distributed mode GridGain nodes, Kafka Connect workers and Kafka brokers are running on different hosts
- GridGain 8.5.3 cluster with GRIDGAIN\_HOME variable set on the nodes
- Kafka 2.0 cluster with KAFKA\_HOME variable set on all brokers

# 1. Run DB Server

We will use H2 Database in this demo.

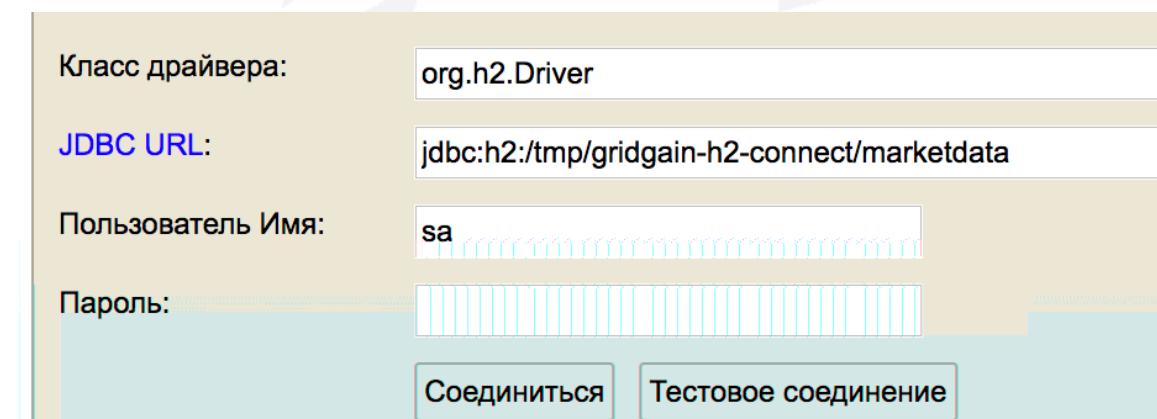
We will use \_\_\_\_\_ as a work directory.

- Download H2 and set H2\_HOME environment variable.
- Run H2 Server:

- In the opened H2 Web Console specify

JDBC URL:

- Press Connect



The screenshot shows the configuration dialog for connecting to an H2 database. It includes fields for the driver class, JDBC URL, username, password, and two buttons at the bottom: 'Соединиться' (Connect) and 'Тестовое соединение' (Test Connection).

Класс драйвера:	org.h2.Driver
JDBC URL:	jdbc:h2:/tmp/gridgain-h2-connect/marketdata
Пользователь Имя:	sa
Пароль:	[REDACTED]
<button>Соединиться</button> <button>Тестовое соединение</button>	

## 2. Create DB Tables and Add Some Data

In H2 Web Console Execute:

- QUOTES
- TRADES
- 
- 
- 

The screenshot shows the H2 Web Console interface. On the left, the database tree view displays the connection URL `jdbc:h2:/tmp/gridgain-h2-connect`, the `QUOTES` table, the `TRADES` table, the `INFORMATION_SCHEMA` schema, users, and the version `H2 1.4.197 (2018-03-18)`. The right side shows two query results. The first result for `TRADES` shows a single row with ID 1 and symbol IBM. The second result for `QUOTES` shows a single row with ID 1, date 2018-10-10, and time 14:47:37.096+100.

```
SELECT * FROM TRADES;
+----+-----+
| ID | SYMBOL |
+----+-----+
| 1  | IBM    |
+----+-----+
(1 строка, 0 ms)

SELECT * FROM QUOTES;
+----+-----+
| ID | DATE_TIME |
+----+-----+
| 1  | 2018-10-10 14:47:37.096+100 |
+----+-----+
(1 строка, 1 ms)
```

### 3. Start GridGain Cluster (Single-node)

```
<bean id="ignite.cfg" class="org.apache.ignite.configuration.IgniteConfiguration">
    <property name="discoverySpi">
        <bean class="org.apache.ignite.spi.discovery.tcp.TcpDiscoverySpi">
            <property name="ipFinder">
                <bean class="org.apache.ignite.spi.discovery.tcp.ipfinder.vm.TcpDiscoveryVmIpFinder">
                    <property name="addresses">
                        <list>
                            <value>127.0.0.1:47500</value>
                        </list>
                    </property>
                </bean>
            </property>
        </bean>
    </property>
</bean>
</property>
</bean>
```

# 4. Deploy Source and Sink Connectors

- Download Confluent JDBC Connector package from  
<https://www.confluent.io/connector/kafka-connect-jdbc/>
- Unzip Confluent JDBC Connector package into
- Copy GridGain Connector package from  
into
- Copy  
into  
property:  
**plugin.path**

Kafka worker configuration file from  
and set the plugin path

# 5. Start Kafka Cluster (Single-broker)

- Configure Zookeeper with

:

- Start Zookeeper:

- Configure Kafka broker: copy default

customize it:

to

- Start Kafka broker:

# 6. Configure Source JDBC Connector

## kafka-connect-h2-source.properties

```
name=h2-marketdata-source
connector.class=io.confluent.connect.jdbc.JdbcSourceConnector
tasks.max=10

connection.url=jdbc:h2:tcp://localhost:19092//tmp/gridgain-h2-connect/marketdata
table.whitelist=quotes,trades

mode=timestamp+incrementing
timestamp.column.name=date_time
incrementing.column.name=id

topic.prefix=h2-
```

# 7. Configure Sink GridGain Connector

## kafka-connect-gridgain-sink.properties

```
name=gridgain-marketdata-sink
topics=h2-QUOTES,h2-TRADES
tasks.max=10
connector.class=org.gridgain.kafka.sink.IgniteSinkConnector

igniteCfg=/tmp/gridgain-h2-connect/ignite-client-sink.xml
topicPrefix=h2-
```

## 8. Start Kafka-Connect Cluster (Single-worker)

./bin/connect-distributed.sh config/connect-distributed.properties

## 9. See Caches Created in GridGain

Open GridGain Web Console Monitoring Dashboard at <https://console.gridgain.com/monitoring/dashboard> and see GridGain Sink Connector created QUOTES and TRADES caches:

Backup	Name	ID	M	On-heap entries				Off-heap entries	
				Total	Primary	Backup	Near	Total	Primary
0	QUOTES	-1895070345	P	0	0	0	0	1	1
0	TRADES	-1812385905	P	0	0	0	0	1	1

# 10. See Initial H2 Data in GridGain

Open GridGain Web Console Queries page and run Scan queries for QUOTES and TRADES:

Key Class	Key	Value Class
hash=-88677777, symbol=IBM, id=11	java.lang.Integer	1

Showing results for scan of TRADES

Key Class	Key	Value Class	Value
java.lang.Integer	1	o.a.i.i.binary.BinaryObjectImpl	TRADES [hash=-88677777, symbol=IBM, id=11]

Showing results for scan of QUOTES

# 11. Update H2 Tables

In H2 Web Console Execute:

- 
- 

The screenshot shows the H2 Web Console interface. The left sidebar lists databases: 'jdbc:h2:/tmp/gridgain-h2-connect' (selected), 'TRADES', and 'INFORMATION\_SCHEMA'. The right pane contains two query results.

Top Query (TRADES):

```
SELECT * FROM TRADES;
```

ID	SYMBOL
1	IBM
2	INTL

(2 строки, 0 ms)

Bottom Query (QUOTES):

```
SELECT * FROM QUOTES;
```

ID	DATE	TIME	PRICE
1	2018-10-10	14:47:37.096	1.0
2	2018-10-10	17:22:31.374	2.0

(2 строки, 0 ms)

# 12. See Realtime H2 Data in GridGain

Open GridGain Web Console Queries page and run Scan queries for QUOTES and TRADES:

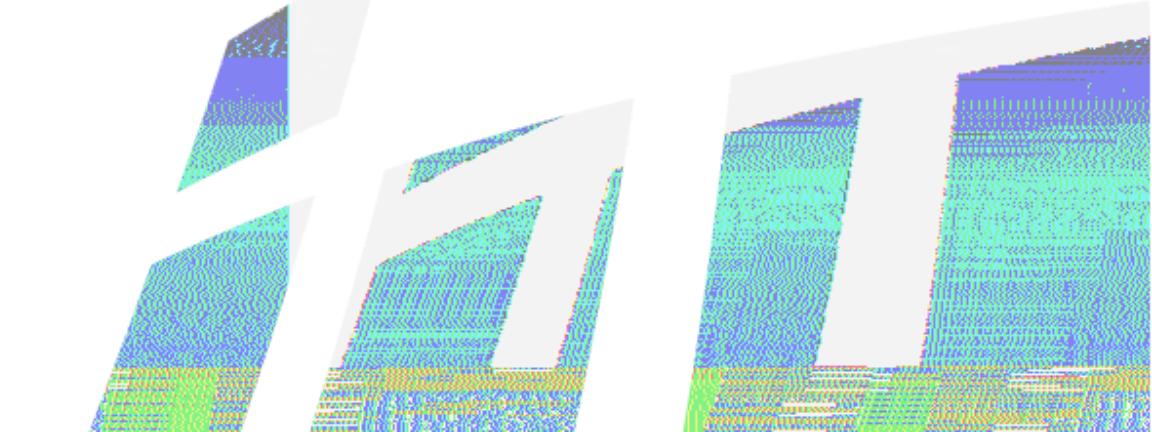
Key Class	Key	Value Class	
java.lang.Integer	1	o.a.i.i.binary.BinaryObjectImpl	TRADES [hash=-88677777, symbol=IBM, id=1]
java.lang.Integer	2	o.a.i.i.binary.BinaryObjectImpl	TRADES [hash=-1299591934, symbol=INTL, id=21]

Showing results for scan of TRADES

Key Class	Key	Value Class	Value
java.lang.Integer	1	o.a.i.i.binary.BinaryObjectImpl	TRADES [hash=-88677777, symbol=IBM, id=1]
4. price=2.0, id=21	java.lang.Integer	o.a.i.i.binary.BinaryObjectImpl	TRADES [hash=-1299591934, symbol=INTL, id=21]

Showing results for scan of QUOTES

# Performance and scalability tuning



# Disable Processing of Updates

For performance reasons, Sink Connector does not support existing cache entry update by default.

Set `processUpdate` configuration setting to `false` to make the Sink Connector update existing entries.

# Disable Dynamic Schema

Source connector caches key and values schemas.

-

# Consider Disabling Schema

Source Connector does not generate schemas if configuration setting is .

Disabling schemas significantly improves performance.

# Carefully Choose Failover Policy

- Can allow losing data? Use **None**.
- Caches are small (e.g. reference data caches)? Use **Full Snapshot**.
- Otherwise use **Backlog**.

# Plan Kafka Connect Backlog Capacity

Only **Backlog** failover policy supports both “at least once” and “exactly once” delivery guarantee.

GridGain Source Connector creates Backlog in the “kafka-connect” memory region, which requires capacity planning to avoid losing data by eviction (unless persistence is enabled).

Consider the worst case scenario:

- Maximum Kafka Connector worker downtime allowed in your system
- Peak traffic

Multiple peak traffic by max downtime to estimate “kafka-connect” data region size.

# Q & A

Thank you!